# BigQuery Cost Report

Project: **demo-warehouse** · Generated 2026-10-15T22:32:44.244616+00:00

Pricing: configured value (buyer's rate) ($6.25 per TiB scanned, on-demand). Read-only; no data leaves this machine.
About this scan: this run processed metadata in 0 read-only query/queries and itself billed $0.00 (typically fractions of a cent, often within the 1 TiB/month free tier). It never reads or changes your table data.
**Note:** This sample was produced from synthetic data, not from a real project - no BigQuery account was contacted. It shows the sections the Community edition produces. The Full edition adds the surprise-bill findings and the partitioning/clustering section.

| Total jobs | Total bytes billed | Estimated query cost |
|---|---|---|
| 27076 | 43,290.61 GB | $264.22 |

Estimated query cost is before the 1 TiB/month free tier and excludes storage.
Result cache: 9,865 of 27,076 jobs (36%) were served from BigQuery's result cache and billed nothing.
Errored jobs: 4 job(s) ended in an error but still billed $0.03 (4.57 GB). Failed queries are charged for what they scanned.

## Cost attribution

### Top query spenders (people)

| Name | Jobs | Bytes | Cost |
|---|---|---|---|
| analyst.one@example.com | 4779 | 1,017.66 GB | $6.21 |
| analyst.three@example.com | 4875 | 994.52 GB | $6.07 |
| analyst.two@example.com | 4780 | 902.16 GB | $5.51 |

### Top service accounts

| Name | Jobs | Bytes | Cost |
|---|---|---|---|
| dbt-runner@demo-warehouse.iam.gserviceaccount.com | 7976 | 38,431.22 GB | $234.57 |
| bi-reader@demo-warehouse.iam.gserviceaccount.com | 4666 | 1,945.05 GB | $11.87 |

### Top datasets/tables

| Name | Jobs | Bytes | Cost |
|---|---|---|---|
| dw_prod.fact_sales | 11566 | 5,769.54 GB | $35.21 |
| dw_prod.fact_events_raw | 210 | 2,621.97 GB | $16.00 |
| dw_prod_stg.stg_sales | 699 | 1,904.72 GB | $11.63 |
| dw_archive.fact_sales_2024 | 2796 | 1,406.32 GB | $8.58 |
| dw_prod_stg.stg_customers | 5918 | 603.82 GB | $3.69 |
| dw_prod.fact_orders | 126 | 437.04 GB | $2.67 |
| dw_prod.dim_customer | 5919 | 297.52 GB | $1.82 |

### Top queries (by cost)

| Query | Jobs | Bytes | Cost |
|---|---|---|---|
| with table_shards_stage as ( select info_tables.table_catalog as table_catalog, info_tables.table_sc | 2638 | 31,656.00 GB | $193.21 |
| select event_type, count(*) as n from `demo-warehouse.dw_prod.fact_events_raw` group by event_type | 210 | 2,621.97 GB | $16.00 |
| select * from `demo-warehouse.dw_prod.fact_sales` where erp_id is not null order by erp_id | 633 | 2,379.50 GB | $14.52 |
| create or replace table `demo-warehouse.dw_prod.fact_sales` as select * from `demo-warehouse.dw_prod | 699 | 1,904.72 GB | $11.63 |
| select erp_id, legal_unit, sum(amount) as total from `demo-warehouse.dw_prod.fact_sales` group by er | 2214 | 1,678.47 GB | $10.24 |

| Query | Jobs | Bytes | Cost |
|---|---|---|---|
| select s.erp_id, a.amount from `demo-warehouse.dw_prod.fact_sales` as s join `demo-warehouse.dw_arch | 2796 | 1,406.32 GB | $8.58 |
| select count(*) as rows_loaded from `demo-warehouse.dw_prod_stg.stg_customers` | 5918 | 603.82 GB | $3.69 |
| select sum(amount) as revenue from `demo-warehouse.dw_prod.fact_orders` | 126 | 437.04 GB | $2.67 |
| select day, sum(amount) as total from `demo-warehouse.dw_prod.fact_sales` where day >= '' group by d | 5919 | 300.68 GB | $1.84 |
| select c.legal_unit, count(*) as customers from `demo-warehouse.dw_prod.dim_customer` as c group by | 5919 | 297.52 GB | $1.82 |

Identical statements are grouped: other queries by their SQL with comments and literal values removed (shortened to fit).

## Beyond this report - the full edition

This community edition shows who and what is driving your BigQuery bill. The full edition adds the analysis that turns that picture into action:

• Surprise-bill risk flags - eight checks covering SELECT * on large tables, queries without LIMIT, missing partition filters, cross-region joins, repeated queries, zombie tables, spend on jobs that failed anyway, and large unpartitioned tables being queried.
• Partitioning and clustering review - which tables are set up to let BigQuery skip data, and which are not.
• Storage cost analysis - what your data costs to store each month, and whether physical (compressed) billing is worth it.
• dbt model attribution - which models are the biggest spenders.
• Daily cost trends and spike detection - which day, and who, caused the jump.
• A no-CLI local web app - pick your project and run the scan from the browser.

For the full edition, email **alex.vinter24@gmail.com**. Read-only, self-hosted, nothing leaves your machine.